

When to reach for which brake

Three axes — **eagerness**, **alignment**, **context rot** — each a brake on a different failure. A field map for when each one fires.

ORDER OF OPERATIONS — a typical sequence; jump in where you need to and skip what doesn't apply

01 · EAGERNESS

/no-eager

discuss the trade-off

02 · ALIGNMENT / PRE

→ **/clarify**

model interviews me

03 · ALIGNMENT / PRE

→ **/align**

model plays back; I confirm

04 · ALIGNMENT / COMMIT

→ **/approval**

plan, then act

05 · ALIGNMENT / COMMIT

→ **/formalize-plan**

heavyweight: written down

01 Eagerness

restrain action

The agent defaults to *doing*. Sometimes there is no action yet — discussion, exploration, a trade-off out loud. Pre-alignment, not alignment.

A TALK ONLY

— no edits should leak

I want to **discuss a trade-off** in design — no edits should land in the repo. > **no-eager**

I'm **thinking out loud** and there may be no target action at all. > **no-eager**

B INFORMATION ONLY

— gather, don't recommend

I'm still exploring — gather **raw material**, not a recommendation. > **no-eager + "research"**

Stress-test the plan — **any gaps**, blind spots, missing inputs? > **no-eager + "gaps"**

- ☐ Discuss a trade-off NO ACT
- ☐ Explore a decision NO ACT
- ☐ Gather context NO ACT
- ☒ Direct request to ship ACT OK

Default failure mode. Without restraint, the agent turns **discussion** into proposals, and proposals into edits. The brake is the conversation shape, not a content filter.

BOUNDARY · DON'T CONFUSE

no-eager — the single brake; tasks that used to be research & gaps are now *no-eager + a noun*.

no-eager vs no-op — *don't act vs. don't summarize*. Different axes; often combined.

02 Alignment

converge before tools fire

Cheap to clarify, expensive to undo. Two sub-modes: *pre-action* (get on the same page) and *at commit* (lock the plan). Different failure modes, different brakes.

A PRE-ACTION

— same page, low confidence

Interview me — pull out what I haven't said. > **clarify**

Read between the lines — what do I want that I **didn't say**? > **latent-demand**

Decompose the work — the **breakdown is the artifact** I'll check intent against. > **breakdown**

B AT COMMIT

— lock the plan, rising confidence

- ☐ Restate *scope & target* so I can correct cheaply
- ☐ Get a *plan written down* before any tool fires
- ☒ Don't run until I *say yes*

Restate scope, target, method, done-when. > **align**

Plan is settled — **write it down** formally before tools fire. > **formalize-plan**

Lightweight gate. **Don't run** until I say yes. > **approval**

Why the split matters. **Pre-action** failure is *working on the wrong thing*. **At-commit** failure is *scope creep mid-execution*. Different brakes for different failures.

BOUNDARY · DON'T CONFUSE

align vs approval — *early framing vs. final gate*. Align checks understanding; approval gates execution.

clarify vs latent-demand — *what I said vs. what I didn't*. Same target, opposite directions.

03 Context rot

defend the window

Models get dumber as the window fills — there's a “dumb zone” (empirically ~100k+ tokens, shifts with each model) where reasoning quality drops. Stay out of it: load cheaply, externalize state, hand off before you cross.

A SUPPRESS NARRATION

— input phase

I'm loading a doc; **read it, don't summarize it** back. > **no-op**

Just absorb this — no recap, no echo, no plan. > **no-op**

B EXTERNALIZE STATE

— working phase

Log decisions, dead-ends, milestones — keep the working set **on disk**, not in the window. > **lab-notebook**

Run a **tight loop** with fresh sub-contexts per story. > **ralph-lnb**

C HAND OFF


— exit phase

Window is filling; **compress this session** into a brief for a fresh agent. > **handoff**

BEFORE · NARRATION CROWDS THE WINDOW



AFTER · SELF-CONTAINED HANDOFF BRIEF



■ loaded materials ■ tool calling ■ user input ■ agent narration
□ headroom

Three phases, three tools. **no-op** on the way in, **lab-notebook** while working, **handoff** on the way out. The **ralph-*** cluster is all three at once — each story gets a fresh window.

BOUNDARY · DON'T CONFUSE

handoff vs lab-notebook — *exit vs. ongoing*. Handoff is one-shot; the notebook persists across the whole project.

no-op vs be-concise — *don't echo vs. shorter echo*. No-op refuses; concise just trims.